

C Programming under Linux

P2T Course, Martinmas 2003–4 *C Lecture 2*

Dr Ralf Kaiser

Room 514, Department of Physics and Astronomy

University of Glasgow

`r.kaiser@physics.gla.ac.uk`

Summary

- Computer Memory Organisation
- Number Systems
- Variables
- Simple Operators
- Printing to the Screen

<http://www.physics.gla.ac.uk/~kaiser/>

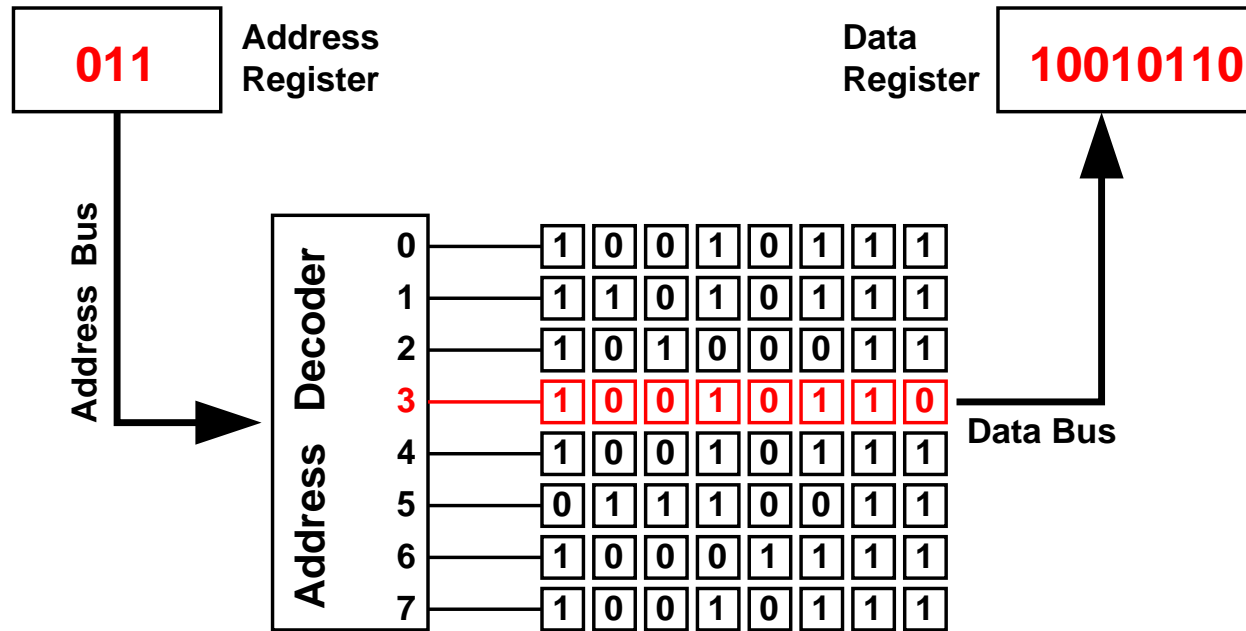
Computer Memory

- Computer memory is classified as either **ROM** (read only memory) or **RAM** (random access memory).
- **Static RAMs (SRAM)** use **flip-flops** as storage elements and can therefore store data indefinitely as long as DC power is applied.
- **Dynamic RAMs (DRAM)** use **capacitors** as storage elements and cannot retain data very long without the capacitors being recharged.
- Data can be read much **faster from SRAMs** than from DRAMs. **DRAMs** can store **much more data** for a given physical size and cost of the memory, because the DRAM cell is much simpler. Therefore you may find that your PC uses DRAM as main memory (e.g. EDO DRAM, extended data out DRAM), but the cache better be SRAM or your PC will be slower than it could be.

Computer Memory cont.

- One **bit** is the unit of information, i.e. one 1 or 0. Eight bits are one **byte**.
- A complete unit of information is called a **word** and generally consists of one or several bytes. As a simplified general rule we can say that memories store data in bytes.
- The memory can be imagined as an array of cells (flip-flops or capacitors) that has 8 columns and a large number of rows. In this picture a 64 MByte memory has 8 columns and 64 million rows (actually , more like 67 million rows, as 1 k is 1024 in memory).
- The location of data in a memory array is called its **address**.

Memory Organisation



- Data units go in and out of the memory on a set of lines called the **data bus**.
- For any read or write operation an address is selected by placing the corresponding binary code on a set of lines called the **address bus**.

Number Systems

- Our regular, every day number system is the **decimal** system. It's probably so successful because of the 10 fingers we have; in fact the word *digit* comes from the Latin word for finger.
- In the decimal system, the position of the digit indicates it's value in powers of 10:
$$2003 = 2 \cdot 10^3 + 0 \cdot 10^2 + 0 \cdot 10^1 + 3 \cdot 10^0$$
- In other number systems the position of the digit indicates it's value in powers of another value: powers of **2** for **binary**, powers of **8** for **octal** and powers of **16** for **hexadecimal** numbers.
- Octal number only need the symbols 0..7, hexadecimal numbers use **A..F** for the numbers **10..15**, e.g. **2003** (decimal) becomes **7D3** (hexadecimal).
- Binary, octal and hexadecimal are the natural number systems for computers.

Variable Types

In C data are stored in variables and there are three basic types of variables:

- `int`
`integer` variables, whose values are `integer numbers` such as 4, -128, 147238
- `float`
`floating point numbers`, corresponding to non-integer numbers with a decimal point. 5.0 is a floating point number, while 5 is an integer.
- `char`
`character` variables with values 'a' to 'z', 'A' to 'Z', '0' to '9' plus punctuation marks, parentheses etc.

Variable Types - int

- size typically reflects the natural size of integers on the host machine. The C standard does not define the size of numbers. Under Linux (or another Unix) integers are typically 32 bits (4 bytes), providing a range from 2147483647 ($2^{31} - 1$) to -2147483648.
- The standard header file `limits.h` defines constants for the various numerical limits.

```
/* Minimum and maximum values a 'signed int' can hold. */  
# define INT_MIN      (-INT_MAX - 1)  
# define INT_MAX      2147483647
```

- for comparison: in Turbo C++ integers are only 16 bit, with a range from 32767 to -32768. The example of 147238 from the last slide would not be possible.

Variable Types - float and double

- usually `float` variables use **4 bytes**, corresponding to a value range of $3.4 \cdot 10^{-38}$ to $3.4 \cdot 10^{+38}$ with a corresponding range of negative numbers and approximately **7 significant digits**.
- `double` variables use **8 bytes**, corresponding to a value range of $1.7 \cdot 10^{-308}$ to $1.7 \cdot 10^{+308}$ with a corresponding range of negative numbers and approximately **15 significant digits**.
- floating point numbers can be given in exponential form, e.g. `1.2e34` is $1.2 \cdot 10^{34}$
- While it is possible to write `.5` instead of `0.5` and `12.` instead of `12.0`, this should be avoided.
- computers do integers better than floating point numbers, and in floating point operations **rounding errors** can accumulate.

Variable Types - char

- single characters are represented by an 8 bit number, this length is also defined in `limits.h`:

```
/* Number of bits in a 'char'. */  
# define CHAR_BIT      8
```

- characters are enclosed in single quotes

```
char char1;    /* first character */  
char char2;    /* second character */  
char char3='z' /* third character, directly initialised */  
  
char2 = '9';   /* assignment statement for second character */
```

- only 7 bits are actively used to encode the character in the **ASCII** code (American Standard Code for Information Interchange), bit 7 can be used as a parity bit.
- characters can also be specified by `\nnn` where `nnn` is their **octal code** according to ASCII (e.g. `\100` is `@`)

Variable Types - Qualifiers

- The variable type can be combined with the qualifiers `short` or `long` to limit or extend their range and with the qualifiers `signed` or `unsigned`.
- where `int` is 32 bit, `short int` is only 16 bit, `long int` is 32 bit, but `long long int` is 64 bit
- `long`, `long long` and `unsigned` numbers are marked with `U` and `L`, `LL` at the end. The largest integer number in `gcc` on a Linux machine is an `unsigned long long int`:

```
/* Maximum value an 'unsigned long long int' can hold.  (Minimum is 0.)  */  
#  define ULLONG_MAX    18446744073709551615ULL
```
- `long double` has a maximum value of `1.18973149535723176502e+4932L` - as defined in `float.h`.

Variable Names

- names in C start with a **letter** or an **underscore** (**_**), followed by any number of letters, numbers or underscores
- for an internal name at least the first **31** characters are significant
- special signs and spaces are not allowed; C commands are **reserved words** and can't be used as variables names
- uppercase is different from lowercase, so **max**, **Max** and **MAX** specify three different variables
- names starting with underscores are conventionally used only for internal and systems variables
- lowercase** variable names are typical for C, so are **lowercase_and_underscore**, but some also use **thisStyle**.

Variable Names - Examples

- examples for valid variable names and declarations are

```
int number_of_students = 47;  /* number of students in this class */
float pi = 3.1415927;         /* pi to 7 decimal places */
int dataRecoil;               /* Recoil detector data */
```

- examples for invalid variable names are

```
3rd_entry    /* starts with a number */
all$done     /* contains a '$' */
int          /* reserved word */
home phone   /* contains a space */
```

- Use variable names that describe the contents. I once knew a programmer who had used the names of rivers in Armenia as variables in a high voltage control program. Beautiful, but not very practical.
- Don't forget to comment the variables, even if the names are chosen well.

Simple Operators

There are 5 simple arithmetic operators in C:

Operator	Meaning
*	multiply
/	divide
+	add
-	subtract
%	modulus (return remainder after integer division)

- multiply (*), divide (\) and modulus (%) have precedence over add (+) and subtract (-)
- parentheses () may be used to group terms
- these operators are also referred to as **binary** operators, because they have two operands, e.g. `a + b`

Simple Operators - Examples

```
int term1;           /* first term */
int term2;           /* second term */
int sum;             /* sum of first and second term */
int difference ;     /* difference of first and second term */
int modulo;          /* term1 modulus term2 */
int product;         /* term1 * term2 */
int ratio ;          /* term1 / term2 */

int main()
{
    term1 = 1 + 2 * 4; /* yields 2*4=8 8+1=9 */
    term2 = (1 + 2) * 4; /* yields 1+2=3 3*4=12 */
    sum = term1 + term2; /* yields 9+12=21 */
    difference = term1 - term2 /* yields 9-12=-3 */
    modulo = term1 % term2 /* yields 9/12=0, remainder is 9 */
    product = term1 * term2 /* yields 9*12=108 */
    ratio = 9 / 12 /* yields 9/12=0 */
    return(sum);
}
```

Floating Point vs Integer Divide

- There is a vast difference between an **integer divide** and a **floating point divide**. In an integer divide the result is **truncated**.
- C allows the assignment of an integer expression to a floating point variable. C will automatically do the conversion. Similarly, a floating point number can be assigned to an integer variable; the value will be truncated.
- The result of a division is **integer** only if both factors are **integers** or if the result is assigned to an **integer variable**. Otherwise it is **floating point**.

Expression	Result	Result Type
19 / 10	1	integer
19.0 / 10	1.9	floating point
19.0 / 10.0	1.9	floating point

Floating Point vs Integer Divide - Example

```

/*****
 * Question:
 *      Why does the following program print:
 *      "The value of 1/3 is 0.0" ?
 *****/
#include <stdio.h>

float answer;    /* The result of our calculation */

int main()
{
    answer = 1/3;
    printf("The value of 1/3 is %f\n", answer);
    return (0);
}

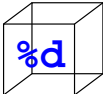
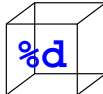
```

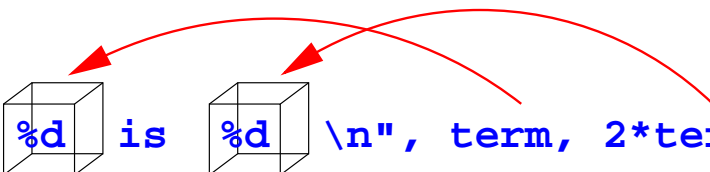
Answer: 1 and 3 are both integers, so the problem is the integer divide where the result is truncated. The expression should be written as

answer = 1.0 / 3.0

The Function `printf`

- The standard function in C used to print something to the screen is `printf`, included in the library `<stdio.h>`.
- The standard form of the `printf` statement is `printf(format, expression-1, expression-2, ...)` where `format` is the string describing what to print, e.g.
- `printf("Twice %d is %d", term, 2*term);` where the special characters `%d` are the **integer conversion specification**. Their place is filled with the values of `expression-1` and `expression-2`; everything else is printed verbatim.

`printf ("Twice  is  \n", term, 2*term);`



- **You** have to check that conversions and expressions match !

Escape Characters

Special characters or **escape characters** starting with '`\`' move the cursor or represent otherwise reserved characters.

Character	Name	Meaning
<code>\b</code>	backspace	move cursor one character to the left
<code>\f</code>	form feed	go to top of new page
<code>\n</code>	newline	go to the next line
<code>\r</code>	return	go to beginning of current line
<code>\a</code>	audible alert	'beep'
<code>\t</code>	tab	advance to next tab stop
<code>\'</code>	apostrophe	character '
<code>\"</code>	double quote	character "
<code>\\</code>	backslash	character \
<code>\nnn</code>		character number nnn (octal)

Format Statements

Besides `%d` there are other conversion specifications, here is an overview of the most important ones:

Conversion	Argument Type	Printed as
<code>%d</code>	integer	decimal number
<code>%f</code>	float	<code>[-]m.dddddd</code> (details below)
<code>%X</code>	integer	hex. number using A..F for 10..15
<code>%c</code>	char	single character
<code>%s</code>	char *	print characters from string until <code>'\0'</code>
<code>%e</code>	float	float in exp. form <code>[-]m.dddddde±xx</code>

In addition, the **precision** and additional spaces can be specified:

`%6d` decimal integer, at least 6 characters wide

`%8.2f` float, at least 8 characters wide, two decimal digits

`%.10s` first 10 characters of a string

Example for printf

Why does $2 + 2 = 1075031184$? (on my laptop; results may vary)

```
#include <stdio.h>

/* Variable for computation results */
int answer;

int main()
{
    answer = 2 + 2;

    printf("The answer is %d\n");
    return (0);
}
```

Answer: The `printf` statement is trying to print a decimal number (%d), but no value is specified - so C makes one up. The proper statement would be

```
printf("The answer is %d\n", answer);
```

ASCII and beyond

- ASCII was defined in 1968 (i.e before most of you were born). It uses **7 bit** and works well with American English, which is after all what it was meant to do.
- ASCII already doesn't work with Danish, French or German - but for this there is an extension to **8 bit** called **Latin1** (aka ISO-8859-1).
- `printf` will typically understand Latin1, so you can use `printf("One beer is 1.95 \xA3");` to print One beer is 1.95 £.
- If you're German (like me) you might want to print 'Schöne Grüße', which you can do like this:
`printf("Sch\366ne Gr\374\337e\n");`
- Of course this still doesn't help you much if you're Thai or Armenian. Have a look at www.unicode.org for Unicode (ISO-10646), a 32 bit code for all the letters in the world.