

C Programming under Linux

P2T Course, Martinmas 2003–4
C Lecture 10

Dr Ralf Kaiser

Room 514, Department of Physics and Astronomy

University of Glasgow

`r.kaiser@physics.gla.ac.uk`

Summary

- Memory Allocation
- Random Numbers
- Bits and Pieces
- Famous Last Words

<http://www.physics.gla.ac.uk/~kaiser/>

Dynamic Memory Allocation

- So far we have always declared all the variables we would need during the running of the program right at the beginning. This can also be referred to as **static memory allocation**.
- What happens if the program finds out that it needs a bit more storage space than we allocated ? Let's say we have an array that we are going to fill as the result of a calculation. What if we don't know from the beginning how long it should be ?
- We can of course allocate just a lot of space, i.e. make the array huge from the beginning. However, this wastes a lot of space and once it's not one array but 100 arrays, we may run out of space.
- The solution lies in **dynamic memory allocation**. The C library contains functions that can allocate storage space at runtime.
- This allows the program to react to demands from the outside and makes it more flexible, e.g. also when it communicates with other programs while running.

Allocating Memory with malloc

- The most basic memory allocation function is `malloc`.
- The function prototype for `malloc` is

```
void *malloc(unsigned int);
```
- The argument of `malloc` is the **number of bytes** to allocate. If `malloc` runs out of memory it returns a null pointer.
- `malloc` returns a pointer to type `void`, i.e. a **generic pointer** that can point to anything, e.g. a string or a structure. What `malloc` doesn't give us is a normal variable with a variable name.
- The returned pointer can be cast to a specific type, but implicit conversion on assignment to a previously defined pointer is allowed in Standard C.
- In C++ the type cast is required.
- In old books you may find that `malloc` returns a pointer to character. That's outdated.

Allocating Memory with malloc - Example

- Suppose we are working on a data base program, using a **structure** called **person**:

```
struct person {  
    char name[30];      /* name of the person      */  
    char address[30];   /* where he/she lives      */  
    int age;            /* how old he/she is       */  
    float height;       /* how tall is he/she in cm */  
}
```

- Instead of using an array of structures **person** we can now use **malloc** to create a **person** when we need one:

```
/* Pointer to a new person structure to be allocated from the heap */  
struct person *new_item_ptr;  
new_item_ptr = malloc(sizeof(struct person));
```

- We don't even have to know the size of a **person** in bytes, the **sizeof** function will figure this out for us.
- The heap is the total of the available memory space, which is large but not infinite. When malloc runs out of space it returns a NULL pointer - and one should check if **malloc** was successful.

De-allocating Memory with `free`

- If we keep allocating memory we will eventually run out. That's why there also has to be a function that de-allocates or frees the memory again. This function is called `free`.
- The function prototype is `void free(void *ptr);`
- In addition the pointer should be set to `NULL`, but doesn't have to. However, it does save us from trying to use freed memory.
- Here is an example using `malloc` to get memory and `free` to dispose of it:

```
const int DATA_SIZE = (16 * 1024) ; /* Number of bytes in the  buffer */
void copy(void)
{
    char *data_ptr;    /* pointer to large data buffer */
    data_ptr = malloc(DATA_SIZE);    /* get the buffer */
    /*
     * use the buffer to do something, i.e. copy a file
     */
    free(data_ptr);
    data_ptr = NULL;
}
```

Other Memory Allocation Functions

- The memory allocation functions are, among others, declared by the header file `stdlib.h`. So if you want to use any of them you have to include it. The other allocation functions are:
 - `calloc(count, size)` allocates a region of memory large enough for an array of `count` elements, each of `size` bytes. The region of memory is declared bitwise to zero.
 - `realloc` changes the size of an already existing pointer while preserving its contents. If necessary, the contents are copied to a new memory region; a pointer to the (possibly new) memory region is returned.
 - `mlalloc` - same as `malloc` but parameter is `unsigned long`
 - `clalloc` - same as `calloc` but parameters are `unsigned long`
 - `cfree` frees memory allocated by `calloc` or `clalloc`.
- For details please have a look at `man` pages, `info` pages, the web or a C book.

Random Numbers

- On occasion you will want to be able to generate **random numbers**, lots of them.
- A typical application in particle and nuclear physics are so-called **Monte Carlo simulations** that use large numbers of simulated physics processes with a statistical distribution of parameters. They are important to model the performance of a detector or to understand the background under the peak that signifies a new particle.
- C offers a function in **stdlib.h** that can be used for this purpose.
- The function prototype is simply

int rand(void);

It returns a pseudo-random integer number between 0 and **RAND_MAX**.

- **RAND_MAX** is defined in **stdlib.h** as 2147483647.

Random Numbers cont.

- The numbers returned by `rand` are pseudo-random, meaning that the distribution is flat, i.e. the probability for any given interval of the same size is about equal.
- However, the sequence is deterministic, i.e. repeated calls to `rand` in one program give different numbers, but if the program writes out 1000 'random' numbers they will always be the same in the same order.
- The function `srand` can be used to change the so-called `seed`, i.e. the input value of the algorithm that works behind `rand`.
- The function prototype here is

```
void srand(unsigned seed);
```

- The same argument in `srand` will produce the same chain of values coming from `rand`. So for 'really' random numbers one can use `rand` to produce a seed for `srand` and then produce a large set of numbers.

Integer and Float Random Numbers - Example

Integer and floating point random numbers (`random.c`).

```
/* integer random numbers */
```

```
f_ptr = fopen("random1.dat", "w");
```

```
for (i=0; i<100000; i++) {  
    fprintf(f_ptr, "%d\n", rand());  
}  
fclose(f_ptr); /* close file */
```

```
/* random numbers between 0 and 1 */  
/* flat distribution */
```



```
f_ptr = fopen("random2.dat", "w");
```

```
for (i=0; i<100000; i++) {  
    number = (float) rand()/RAND_MAX;  
    fprintf(f_ptr, "%f\n", number);  
}  
fclose(f_ptr); /* close file */
```






- The standard output of `rand` consists of integers. However, you may want to have floating point values between 0 and 1 instead.
- This can be constructed in a simple fashion (see example).
- `stdlib.h` also contains additional functions that do this for you, but `rand` is Standard C.

Integer and Float Random Numbers - Example

Integer and floating point random numbers (`random.c`), histograms of random number output.



Gaussian Random Numbers



- The key to a **Gaussian distribution** is that it can be created from a flat distribution by **sampling**.
- Already the **sum of 6** random numbers from a flat distribution is very close to a **Gaussian**.


Gaussian Random Numbers - Example

Gaussian random numbers (`random.c`), histogram of random number output with Gaussian fit.

```
/* random numbers between 0 and 1 */
/* Gaussian distribution around 0.5 */

f_ptr = fopen("random3.dat", "w");

for (i=0; i<100000; i++) {
    number = 0.0;
    for (j=0; j<6; j++) {
        number += (float) rand()/RAND_MAX;
    }
    fprintf(f_ptr, "%f\n", number/6.0);
}
fclose(f_ptr); /* close file */
```



(The histograms on this slide and the previous one were produced using PAW, a data analysis package from CERN which is available for Linux, but unfortunately based on FORTRAN. The successor is called root and based on C++.)

The Ternary Operator '?:'

- We have seen **unary** operators (like '++') that act on one variable and **binary** operators (like '+') that require two.
- There is also a single **ternary** operator, the operator '?:'. It's use is similar to an **if/then/else** construction, but in contrast to **if/then/else** '?:' can be used inside of an expression.

- The general form is

(expression) ? value1 : value2

- For example, the following construct assigns to **amount_owed** the value of the balance or zero, depending on the amount of the balance:

amount_owed = (balance < 0) ? 0 : balance;

- The following macro returns the minimum of it's two arguments:

#define min(x,y) ((x) < (y) ? (x) : (y))

The do/while Statement

- The `do/while` statement has the following syntax:

```
do {  
    statement;  
    statement;  
} while (expression);
```

- The program will loop, test the expression, and stop if the expression is false (0).
- This construct will always **execute at least once**.
- It's not frequently used in C programming. Most programmers prefer to use a `while/break` combination.

The Comma (,) Operator

- The **comma operator** (,) can be used to group statements, for example instead of

```
if (total < 0) {  
    printf("You owe nothing\n");  
    total = 0;  
}
```

we can also write

```
if (total < 0)  
    printf("You owe nothing\n"), total = 0;
```

- In most cases curly braces {} should be used instead. The only places where the comma operator is useful is for the declaration of a couple of simple variables in one line

```
int i, j, k; /* a couple of counters */
```

and in **for** statements:

```
for (two = 0, three = 0;  
     two < 10;  
     two += 2, three += 3)  
    printf("%d %d\n", two, three);
```


Finally: The `goto` Statement

- You thought perhaps that `goto` has been banished to forever live only in damp dungeons of BASIC. You were wrong. It exists, but you don't every have to use it.
- For those infrequent times when you actually want to use a `goto`, the correct syntax is

`goto label;`

where `label` is a statement label.

- Example:

```
for (x = 0; x < X_LIMIT; x++) {  
    for (y = 0; y < Y_LIMIT; y++) {  
        if (data[x][y] == 0)  
            goto found;  
    }  
}  
printf("Not found\n");  
exit(8);
```

```
found:  
    printf("Found at (%d,%d)\n", x, y);
```

Things not to do

- '=' is not the same as '==' ! Don't mix them up !
- C starts counting at '0' not at '1' ! Stick to this, don't start at 1 !
- `matrix[10,12]` is not the correct C syntax !
- Don't forget the `\0` at the end of a string !
- Don't use 8 spaces at the begin of a line in a Makefile instead of a `tab` !
- Never put an assignment inside another statement !
- Don't make a block of code inside a function longer than a few (about 3) pages !

Things to do

- Use `++` and `--` on lines by themselves.
- One variable declaration per line.
- Use comments. Lots of them. Useful ones.
- Make variable names explicit. 'total' is better than 't'.
- Always put a `default: case` into a `switch` statement.
- Use `const` instead of `#define` wherever possible.
- Put `()` around each argument of a preprocessor macro.
- Generally use enough `()` if you don't know the precedence rules. (And you haven't learned them here...).
- Use a Makefile.
- Make a backup of your work. Then make another one. If it's important make one more. For my thesis I made one on another disk, one on another machine, one on another continent and one on tape - every evening.